



# Facial Expression Recognition

By: Ethan Assefa, Thomas Burrell, Tatev Gomtsyan

# AGENDA

1

## MOTIVATION

FER uses and goals

2

## OUR DATA

Description of data

3

## MODELS

Development process

4

## CONCLUSION

Model results and comparison

# MOTIVATION

Facial Emotion Recognition (FER) helps create more intuitive and human-like interactions between computers and people, leading to more effective and engaging communication.

Cross-industry applications (healthcare, security, marketing, etc.)

- Virtual assistants
- Customer service bots
- Social robots

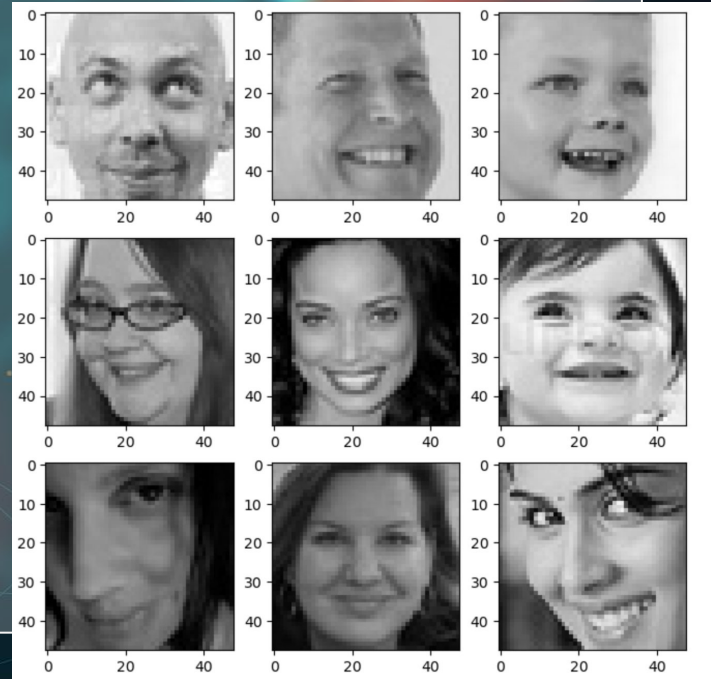
Deep Learning allows for development of complex models that can recognize a wide range of expressions with high accuracy. CNNs are powerful pattern recognition tools in images. We explored a variety of techniques and models to discover the best approach to FER.

# OUR DATA

Kaggle 'face expression recognition' dataset:

- ~29,000 images in Train folder
- ~7,000 images in Validation folder
- Seven expression categories

- angry
- disgust
- fear
- happy
- neutral
- sad
- surprise



# APPROACH: MODEL FROM SCRATCH

Model building	Parameters	Output
<ul style="list-style-type: none"><li>• Image preprocessing</li><li>• 4 CNN Layers, flattened</li><li>• 2 fully connected layers</li><li>• Initialize Adam optimizer</li></ul>	<ul style="list-style-type: none"><li>• Learning rate: 0.001 (Adam adaptively adjusts based on gradients and previous updates during training)</li><li>• Epochs: 20</li><li>• Batch size: 64</li></ul>	<ul style="list-style-type: none"><li>• Updated training and validation accuracies</li><li>• Loss Value</li><li>• Precision and Recall scores</li><li>• Testing model on new images</li></ul>



# TensorFlow/Keras: Model from Scratch

```
no_of_classes=7

model=Sequential()

#1st CNN layer
model.add(Conv2D(64,(3,3),padding="same",input_shape=(48,48,1)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

#2nd CNN layer
model.add(Conv2D(128,(5,5),padding="same"))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

#3rd CNN layer
model.add(Conv2D(512,(3,3),padding="same"))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

#4th CNN layer
model.add(Conv2D(512,(3,3),padding="same"))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
```

## FEATURES

- 4 CNN layers, flattened
- Batch Normalization
- Relu activation
- Max pooling
- Dropout
- 2 fully connected layers

Epoch 1

val\_accuracy: 0.4037

Epoch 1/10

/Users/tatevgomtsyan/anaconda3/lib/python3.11/site-packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:120: UserWarning: Your `PyDatasetAdapter` could call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.

225/225 — 851s 4s/step — accuracy: 0.2561 — loss: 1.9544 — val\_accuracy: 0.4037 — val\_loss: 1.6660 — learning\_rate: 0.0010

Epoch 9

val\_accuracy: 0.5416

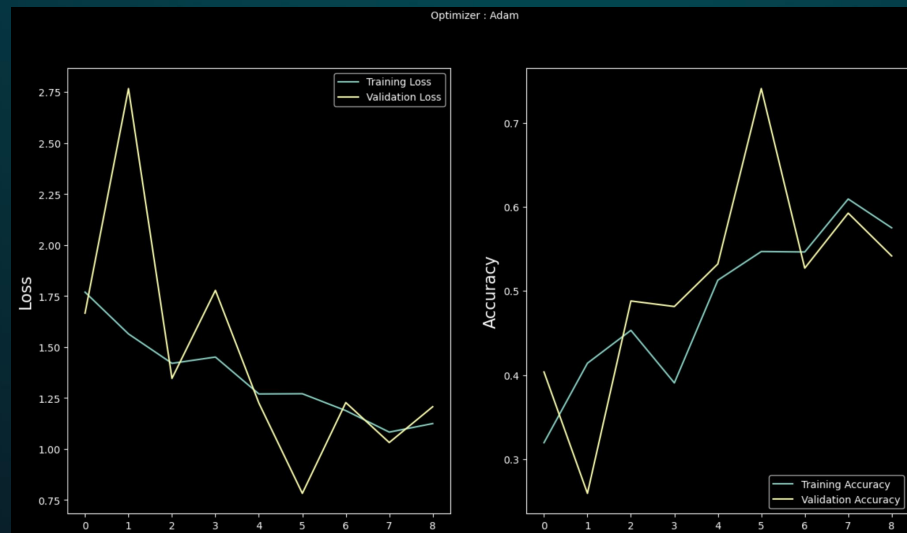
225/225 — 0s 3s/step — accuracy: 0.5790 — loss: 1.1147

Epoch 9: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

225/225 — 757s 3s/step — accuracy: 0.5790 — loss: 1.1148 — val\_accuracy: 0.5416 — val\_loss: 1.2074 — learning\_rate: 0.0010

Epoch 9: early stopping

Restoring model weights from the end of the best epoch: 6.



# PyTorch: Pretrained Model

```
class FaceModel12(nn.Module):
    def __init__(self):
        super(FaceModel12, self).__init__()
        self.eff_net = timm.create_model('resnet101', pretrained=True, num_classes=7)
        self.dropout = nn.Dropout(0.5) # Dropout Layer with 50% probability
        self.batch_norm = nn.BatchNorm1d(7) # BatchNorm for 7 classes output from the model

    def forward(self, images, labels=None):
        logits = self.eff_net(images)
        logits = self.dropout(logits)
        logits = self.batch_norm(logits)

        if labels is not None:
            loss = nn.CrossEntropyLoss()(logits, labels)
            return logits, loss
        return logits
```

## FEATURES

- Using pre-trained model from 'timm' (Torch Image Models) library
- Create model named FaceModel()
- ResNet-101 backbone
- pretrained=True (initialized weights from large dataset)

## FINE-TUNING

- Data augmentation
- Epochs = 20
- Learn rate = 0.001
- Batch Size = 64

Epoch 1

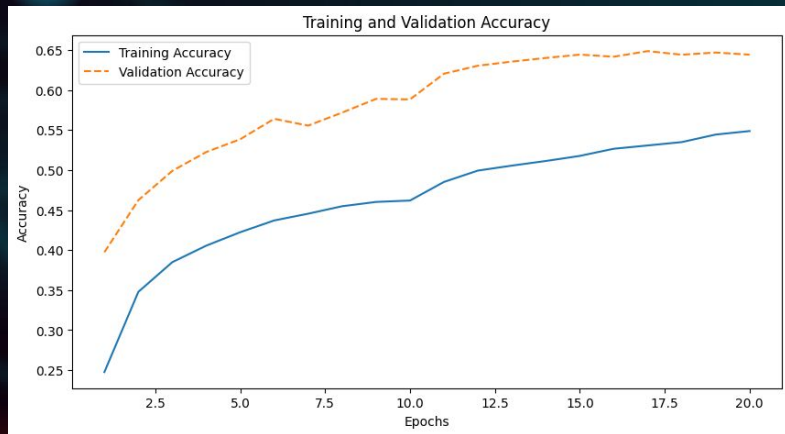
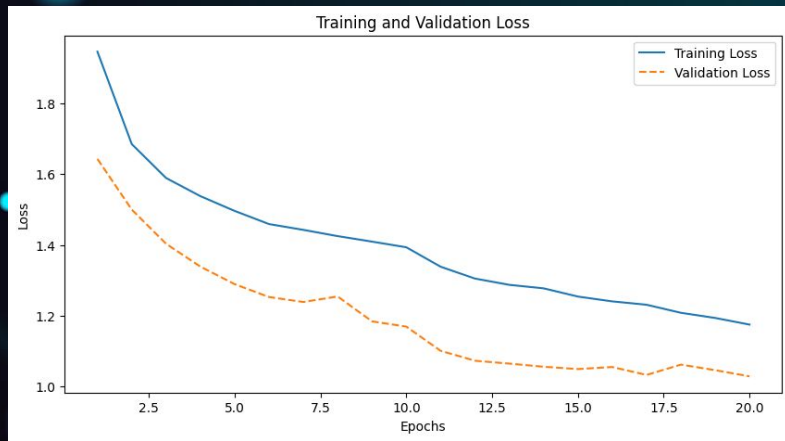
```
EPOCHS[TRAIN]1/20: 100%|██████████| 451/451 [27:36<00:00, 3.67s/it, loss=1.493357, acc=0.415629, precision=0.415308, recall=0.415308]
EPOCH[VALID]1/20: 100%|██████████| 111/111 [02:43<00:00, 1.48s/it, loss=1.235303, acc=0.528721, precision=0.527805, recall=0.527805]
Saved Best Valid Loss
Saved Best Precision
Saved Best Recall
Epoch 1:
Train - Loss: 1.493357, Acc: 0.415629, Precision: 0.415308, Recall: 0.415308
Valid - Loss: 1.235303, Acc: 0.528721, Precision: 0.527805, Recall: 0.527805
```

Epoch 20

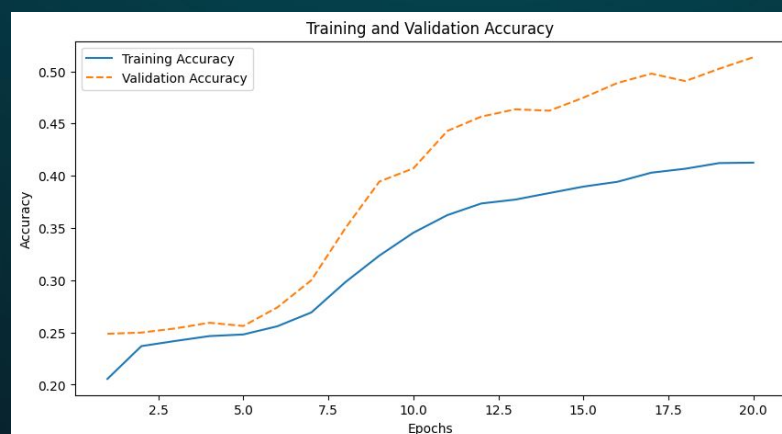
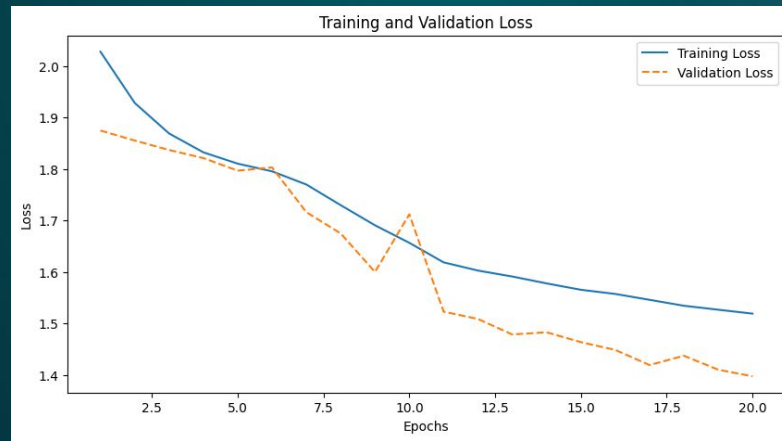
```
EPOCH[TRAIN]20/20: 100%|██████████| 451/451 [00:20<00:00, 22.37it/s, acc=tensor(0.5489), loss=1.17]
EPOCH[VALID]20/20: 100%|██████████| 111/111 [00:02<00:00, 49.92it/s, acc=tensor(0.6443), loss=1.03]
Saved Best Model - Weights with Lowest Validation Loss
Epoch 20/20: Train Loss: 1.1747, Accuracy: 0.5489, Precision: 0.5526, Recall: 0.5526
Validation Loss: 1.0284, Accuracy: 0.6443, Precision: 0.5659, Recall: 0.5659
```

Accuracy: 0.6443,

# ResNet101 Results



# EfficientNet (B5) Results





# MODEL COMPARISON

Model Type	Accuracy
Individual Emotion Model	Angry - 86.41%    Disgust - 98.43% Fear - 85.59%    Happy - 74.17% Neutral - 82.79%    Sad - 83.88% Surprise - 88.72%
Multi-class General Model (pyTorch):	42.15%
Ensemble Model (Single-Emotion and Multi-Class Models):	45.08%
Ensemble Model (Single-Emotion Models Only):	40.16%
Stacking:	40.83%
<b>Resnet101 (WINNER!)</b>	<b>64.43%</b>
EfficientNet (B5):	51.34%

# MODEL TYPE PROS/CONS

Pre-trained	Model from Scratch
<ul style="list-style-type: none"><li>● Speeds up Development</li><li>● Requires Less Data</li><li>● Generalization Over Specialization</li><li>● Limited Customization</li></ul>	<ul style="list-style-type: none"><li>● Highly customizable</li><li>● Optimized for Specific Data</li><li>● Requires Large Datasets</li><li>● Time and Resource Intensive</li></ul>

# CONCLUSION

## CHALLENGES

- Variability in emotional expressions, hard to create model that generalizes across diverse populations
- Cross-cultural differences, would need a more diverse dataset to ensure accuracy and inclusivity
- Privacy and ethical concerns, consent, potential misuse

## CONSIDERATIONS

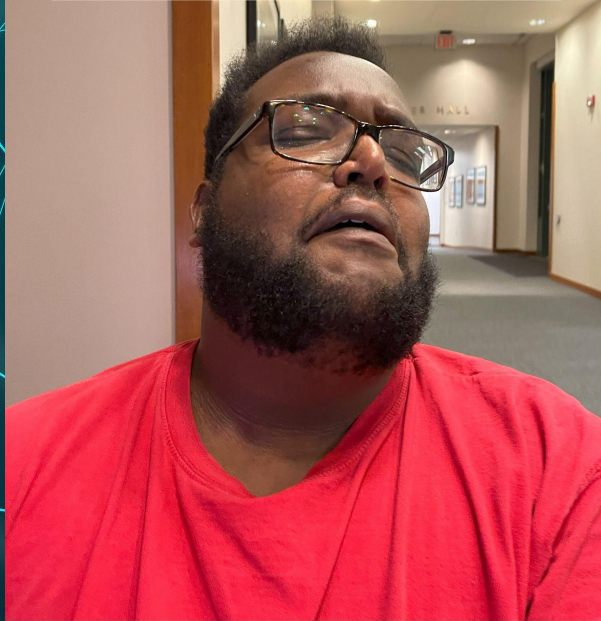
- Make our own train/validation/test split from data
- Further fine-tuning parameters (more epochs, batch size)
- Better computational infrastructure (more resources)

## FUTURE WORK:

- Real-time live video processing for mental health interventions (during telehealth or therapy sessions, with consent)
- Surveillance and security



**Actual: Disgusted**  
**Predicted: Disgust**



**Actual: Surprised**  
**Predicted: Disgust**



**Actual: Happy**  
**Predicted: Happy**



ResNet Model: Image tatev\_surprised.jpg is classified as Disgust  
ResNet Model: Image ethan\_disgust.jpeg is classified as Disgust  
ResNet Model: Image trey\_smile.jpeg is classified as Happy



**THANK YOU!**  
**Questions?**